

# CS2102 Cheatsheet

## General

A *key* is a minimal *superkey*, i.e. no proper subset of itself is a superkey

- A *superkey* is a subset of attributes that unique identifies its tuples
- Key attribute values cannot be null (*key constraints*)
- Can have multiple keys (*candidate keys*). One is the *primary key*.

A *foreign key* refers to primary key of a second relation (can be itself)

- Each foreign key value must be the primary key value in the referenced relation or be null (*foreign key constraint*)

## SQL Data Types

boolean, integer, float8, numeric, numeric(p, s), char(n), varchar(n), text, date, timestamp

## Create/Drop Table

```
create table Students (  
    dept          varchar(20) default 'CS'  
);  
drop table if exists Students cascade;
```

Is Null Predicate: `x IS NULL`

Is Distinct From Predicate: `x IS DISTINCT FROM y`

Equivalent to `x <> y` if both `x` and `y` are non-null, else is false if both are null, else true if only one is null

**not null Constraint:** `name varchar(100) not null,`

**unique Constraint:** `studentId integer unique, or unique (city, state) -- at bottom`

**primary key Constraint:** `studentId integer primary key, or primary key (sid, cid) -- at bottom`

**foreign key Constraint:** `studentId integer references Student (id), or foreign key (a, b) references Other (a, b) -- at bottom`

If any of the referencing columns is null, a referencing row can escape satisfying the FK constraint. If we add **match full** to the end of constraint, then a referencing row only escapes if all of the referencing columns are null.

**check Constraints**

```
check (day in (1,2,3,4,5)), -- besides day  
check ((hour >= 8) and (hour <= 17)), -- at bottom
```

## Constraint Names

```
bDate      date constraint bdate check (bdate is not null) or  
constraint obj_pri_key primary key (name, day, hour)  
Insert: insert into Students (name, studentId) values ('Bob',  
67890), ('Carol', 11122);  
Delete: delete from Students where dept = 'Maths';  
Update: update Accounts set balance = balance + 500, name =  
'Alice' where accountId = 12345;
```

## Foreign Key Constraints Violations

*No Action:* Rejects action if it violates constraint (default). *Restrict:* Same as No Action but constraint checking is not deferred. *Cascade:* Propagates delete/update to referencing tuples. *Set Null:* Updates foreign keys to null. *Set Default:* Updates foreign keys to some default value. We will need to specify this value at the referencing column, and it must meet the foreign key constraints, else action will fail

**foreign key (a, b) references Other (a, b) on delete cascade**

**Transactions**

Starts with **begin;** and ends with **commit;** or **rollback;**. Can contain multiple SQL statements.

*Atomicity:* Either all effects are reflected or none. *Consistency:* Executed in

isolation, preserves DB consistency. *Isolation:* From the effects of other concurrent transactions. *Durability:* Effects persist even if system fails.

## Deferrable Constraints

unique, primary key and foreign key constraints can be deferred using **deferrable initially deferred** or **deferrable initially immediate**

We can change this using **set constraints**, e.g. **set constraints fkey deferred;** or **set constraints fkey immediate;** (retroactive)

## Modifying Schema

```
alter table Students alter column dept drop default;  
alter table Students drop column dept;  
alter table Students add column faculty varchar(20);  
alter table Students add constraint fk_grade foreign key  
(grade) references Grades;  
Expressions in Select Clause
```

- **distinct:** Removes duplicates
- **A as B:** Rename output column A to B, column alias
- **'String' || value:** String concatenation
- **round(A):** Rounds to nearest integer value

## Set Operations

**union, intersect** and **except** are same as  $\cup, \cap, -$ . Eliminates duplicates.

**union all, intersect all** and **except all** preserve duplicate records.

Example: `select A from R except select B from S;`

## Multi-Relation Queries

`Rfull as R, Sfull as S` or `Rfull R, Sfull S`: Cross-product

`R inner join S on R.A = S.A, R join S on R.A = S.A`

`R natural join S`

`R left outer join S on R.A = S.A, R left join S on R.A = S.A`

`R natural left outer join S, R natural left join S`

## Subqueries

**exists:** Returns true if subquery is non-empty. E.g. **where exists**

**in:** Subquery must return a single column, and true if it is non-empty and one row equals expression. E.g. **where expression in**

**any:** Same as **in** but instead of equality, checks for result of binary operation.

E.g. **where price > any**

**all:** Same as **any** but requires all to match, e.g. **where price >= all**

- To use the above with multiple columns, use row constructors, e.g. **where row(day, hour) <= all** or **where (day, hour) <= all**.
- Can be used in **where, from** and **having**.
- If used in **from**, it must be enclosed in parentheses and assigned table alias. E.g. `select * from (subquery) as persons(name, age);`

## Scalar Subqueries

A subquery that returns at most one tuple with one column. If empty, then null is returned. It can be then used as a scalar expression. Also called **nested queries** or an **inner query** nested within an **outer query**.

In terms of table alias (tuple variable) scoping rules, it's the same as a normal language. If the inner query refers to a tuple variable declared in an outer query, then it is a **correlated nested query**.

**Order By:** `order by area asc, price desc; order by area, price desc;`

**Limit:** `limit 3:` Show top 3.

**Offset:** `offset 3:` Show 4<sup>th</sup> and onwards.

## Aggregate Functions

`min(A), max(A), avg(A), sum(A), count(A)` (non-null), `count(*)` (all rows, may be null), `avg(distinct A)` (non-null), `sum(distinct A)` (non-null), `count(distinct A)` (non-null).

- Performed on an empty relation OR a column entirely of **nulls**, all functions except **count** will return **null**. **count** will return **0**, unless **count(\*)** is used, in which case **n** will be returned.
- Can be used in **select, having** and **order by**.
- If a **select** clause contains an aggregated function and there is no **group by** clause, then the **select** clause must not contain any column that is not in an aggregated expression, i.e. all or nothing.

## Group By

Conditions for a **group by** clause on relation R:

- Output column **A** must appear in the **group by** clause, OR
- **A** appears in an aggregated expression in the **select** clause.
- The primary key of R appears in the **group by** clause.

## Having

This is similar to **where**, but is applied after **group by**.

Conditions for a **having** clause on relation R:

- Column **A** that appears in the **having** clause must appear in the **group by** clause, OR
- **A** appears in an aggregated expression in the **having** clause.
- The primary key of R appears in the **group by** clause.

## Conceptual Evaluation of Queries

<b>select</b>	<b>distinct select-list</b>
<b>from</b>	<b>from-list</b>
<b>where</b>	<b>where-condition</b>
<b>group by</b>	<b>groupby-list</b>
<b>having</b>	<b>having-condition</b>
<b>order by</b>	<b>orderby-list</b>
<b>offset</b>	<b>offset-specification</b>
<b>limit</b>	<b>limit-specification</b>

1. Compute cross-product of tables in from-list
2. Select tuples that evaluate to true for the where-condition
3. Partition selected tuples into groups using groupby-list
4. Select the groups that evaluate to true for the having-condition
5. For each selected group, generate an output tuple based on select-list
6. Remove duplicate output tuples because of distinct
7. Sort the output tuples based on orderby-list
8. Remove the appropriate output tuples based on offset-specification and limit-specification

## Common Table Expressions

We can declare common table expressions shared by all subsequent queries.

```
with  
R1 as (Q1),      -- comma for non-last CTEs  
R2 as (Q2)       -- no comma for final CTE
```

## Views

Virtual relation that can be used for querying.

**create or replace view v3 (rname1, rname2) as**

This creates an interface, where users/applications query using the schema of the view, thus safe from changes to the actual logical/physical schema.

## Case

```
case  
    when marks >= 70 then 'A'  
    else 'D'  
end as grade  
case grade          -- case expression  
    when 'A' then 70  
    else 0  
end as marks
```

**Coalesce:** select coalesce(third, second, first) as result  
Returns the first non-null value in its arguments.

**Nullif:** select name, nullif(result, 'absent') as status  
Returns null if result is equal to 'absent', otherwise result.

**Like / Similar To:** where cname like '\_\_\_%e';  
Underscore matches a single character

- % matches any sequence of 0 or more characters
- More complex regex will need similar to.

**Functions in SQL**  
create or replace function a(x int) returns char(1) as \$\$  
-- select a single character using x  
\$\$ language sql;  
returns R  
returns setof R  
d(out x int, out y int) returns record as \$\$ -- called using  
SELECT d(); returns single tuple  
e(out x int, out y int) returns setof record as \$\$  
f() returns table(x int, y int) as \$\$ -- same as e  
You can use inout for the parameters if the input “schema” and the output “schema” share columns.

**Procedures**  
create or replace procedure g(x int) as \$\$ -- CALL g(x);

**Variables and Control Structures**

```
as $$                                if condition1 then
declare                             statement1;
temp_val integer;                   elsif condition2 then
begin                               statement2;
-- function body                     else
end;                                else-statement;
$$ language plpgsql;                end if;
```

```
loop
exit when condition;
statements;
end loop;
```

**Cursor**  
declare  
curs cursor for (select \* from R order by A desc);  
r record;  
begin  
open curs;  
loop  
fetch curs into r;  
exit when not found;  
-- assign output table column names to values  
return next;  
end loop;  
close curs;

end;  
fetch prior from cur into r; fetch first from cur into r;  
fetch last from cur into r; fetch absolute 3 from cur into r;

**Triggers**  
create or replace function func() returns trigger as \$\$  
create trigger trigger\_name [before / after]  
[insert / update /delete] on R for each [row / statement]  
execute function func();  
Order of triggers: before statement, before row, after row, after statement

**Special Values in Triggers**

**TG\_OP:** Operation that activates the trigger, i.e. 'INSERT', 'UPDATE', 'DELETE', 'TRUNCATE'

**TG\_TABLE\_NAME:** Name of table causing the invocation

**OLD:** Old tuple being updated or deleted. null for insertion.

**NEW:** New tuple being inserted or updated. null for deletion.

**Instead Of Triggers (Only Row-Level)**  
You can also define instead of [insert / update / delete] triggers for views (only), e.g. to update the actual tables instead.

**Return Values for Row-Level Triggers**  
before insert: non-null t – t will be inserted, null – no insertion  
before update: non-null t – t will be the updated tuple, null – no update  
before delete: non-null t – deletion happens, null – no deletion  
after [insert/update/delete]: return value does not matter  
instead of: non-null t – proceed, null – ignore operations on current row  
If a before row-level trigger returns null, then all subsequent triggers on the same row are omitted.

**Return Values for Statement-Level Triggers**  
Return values are ignored. Raise exceptions if to ignore operations.

**Trigger Condition**  
Example: for each row when (NEW.Name = 'Hello') execute  
No select in when, no OLD in when for insert, no NEW in when for delete, no when for instead of.

**Deferred Triggers**  
create constraint trigger trigger\_name after ... on R deferrable initially [deferred / immediate] for each row ...  
Only works for after and for each row.

If initially immediate, we need to change on the fly  
begin transaction; set constraints trigger\_name deferred;

**Functional Dependencies**  
A → B means A decides B, i.e. if two rows have the same A, then they have the same B.

Armstrong’s Axioms

- Reflexivity: ABC → A (set to subset)
- Augmentation: If A → B, then AC → BC for any C
- Transitivity: If A → B and B → C then A → C

Extended Axioms

- Decomposition: If A → BC then A → B and A → C
- Union: If A → B and A → C then A → BC

**Closures**  
{A}<sup>+</sup> denotes the set of attributes that can be directly or indirectly decided by A, also called the closure of A.  
To compute: start with {A}; for all FDs such that the LHS can be found in the current closure, put the RHS into the closure as well; repeat until no more new attributes can be added.  
To prove X → Y, we just need to show {X}<sup>+</sup> contains Y. Opposite is true.

**Keys, Superkeys and Prime Attributes**  
Superkey: A set of attributes in a table that decides all other attributes. Key: A superkey that is minimal. Prime attribute: Attribute that appears in a key.  
To find keys of T: consider all subsets of T; derive the closure of each subset; identify the superkeys; identify the keys.

Heuristic: Check smaller attribute sets first.

Heuristic: If an attribute does not appear in the RHS of any FD, then it must be in every key.

**Non-Trivial and Decomposed FDs (NTD)**

*Decomposed FD:* RHS only has one attribute. *Non-trivial and decomposed:* RHS does not appear in LHS.

To find such FDs, simply compute all closures, remove trivial attributes, then separate them into decomposed FDs.

**Boyce-Codd Normal Form (BCNF)**  
*Normal form:* Definition of minimum requirements in terms of redundancy.  
A table R is in BCNF if every NTD has a **superkey** as its LHS.  
To check if R is in BCNF: compute all closures; find the keys from the closures; derive NTD FDs from the closures; check the BCNF requirement.

Heuristic: Check if there’s a closure that contains “more but not all”, i.e. the RHS has more attributes than the LHS, but does not contain all attributes.

**BCNF Decomposition**  
Find any “more but not all” closure {X}<sup>+</sup>, then decompose R into R<sub>1</sub> and R<sub>2</sub>:

- R<sub>1</sub> contains all attributes in {X}<sup>+</sup> (i.e. RHS) and
- R<sub>2</sub> contains all attributes in X and attributes not in {X}<sup>+</sup> (i.e. LHS + remaining attributes).

Then, project the closures from R onto R<sub>1</sub> or R<sub>2</sub>, then repeat the process.  
To project: enumerate all attribute subsets in R<sub>1</sub> (WLOG); derive their closures on R; remove irrelevant attributes.

If a table only has two attributes, then it must be in BCNF.

**BCNF Properties**  
No update or deletion or insertion anomalies, small redundancy and you can always reconstruct the original table.  
A lossless join is guaranteed whenever the common attributes in R<sub>1</sub> and R<sub>2</sub> constitute a superkey of R<sub>1</sub> or R<sub>2</sub>.  
But it may not guarantee **dependency preservation**.

**Dependency Preservation**  
Let S be the given set of FDs on the original table, and S' be the set of FDs on the decomposed tables.  
A decomposition preserves all FDs iff S and S' are equivalent, i.e. every FD in S' can be derived from S, and vice versa.  
Preserving FDs allow us to prevent inappropriate updates.

**Third Normal Form (3NF)**  
A table satisfies 3NF iff for every NTD, either the LHS is a superkey or the RHS is a prime attribute.  
This is more lenient than BCNF, i.e. satisfying BCNF → satisfying 3NF.  
The checking is similar to BCNF, except that the requirement is different.

**3NF Decomposition**  
3NF decomposition does a single split into two or more parts.  
Derive the minimal basis of the FDs; combine the FDs whose LHS are the same, i.e. A → B and A → C becomes A → BC; create a table for each remaining FD, e.g. R1(A, B, C); if none of the tables contain a key of the original table R, create a table that contains any key of R.

**Minimal Basis / Cover (MB)**  
We simplify a set S of FDs using the four conditions:

1. Every FD in the MB can be derived from S and vice versa.
2. Every FD in the MB is a NTD.
3. No FD is redundant, i.e. can be derived from other FDs in the MB.
4. For each FD, none of the LHS attributes is redundant, i.e. if we remove it, it cannot be derived from the original set of FDs.

Algorithm for MB  
Transform the FDs so that each RHS only contains one attribute; remove redundant attributes on the LHS of each FD; remove redundant FDs.