

1. This question considers a binary relational algebra operator called the **division operator** denoted by $/$ ¹.

Consider two relations R and S where the set of attributes in the schema of R and S are $(A_1, \dots, A_m, B_1, \dots, B_n)$ and (B_1, \dots, B_n) respectively where $m \geq 1$ and $n \geq 1$. That is, the set of attributes in S is a *proper* subset of the set of attributes in R .

Assume that the attributes that are in R but not in S are ordered as (A_1, \dots, A_m) in the schema of R and the schema of S is (B_1, \dots, B_n) . Let L denote the list of attributes in the schema of R .

The division of R by S (denoted by R/S) computes the largest set of tuples $Q \subseteq \pi_{A_1, \dots, A_m}(R)$ such that for every tuple $(a_1, \dots, a_m) \in Q$,

$$\pi_L(\{(a_1, \dots, a_m)\} \times S) \subseteq R$$

Q is also referred to as the **quotient** of R/S and its schema is (A_1, \dots, A_m) . The following example illustrates R/S given two relations $R(A, B)$ and $S(B)$.

R			
A	B		
a	1		
a	2		
b	1		
c	1		
c	2		
c	3		
d	2		
d	3		

S		R/S	
B		A	
1		a	
2		c	

- (a) Consider again Question 5 in Tutorial 1 to find the restaurants that sell all the pizzas that Maggie likes and don't sell any pizza that Ralph likes. Write a relational algebra expression for this query that uses the division and natural join operators.
- (b) Given relations $R(A, B)$ and $S(B)$, write a relational algebra expression to compute the division of R by S using only the basic relational operators (*i.e.*, σ , π , ρ , \times , \cup , \cap and $-$).

Solution: If you noticed, R/S given $R(A, B)$ and $S(B)$ is exactly all the tuples with the same value on attribute A for which the attribute B contains all values in S (*i.e.*, the superset of S).

- (a) In this case, finding the restaurant that sells all the pizzas that Maggie likes can be done using the division operator.

$$Q_1 := \pi_{\text{rname, pizza}}(\text{Sells}) / \pi_{\text{pizza}}(\sigma_{\text{cname}=\text{'Maggie'}}(\text{Likes}))$$

¹It is a "division" operation because it kind of reverses the "multiplication" operation \times .

We then need to remove all the restaurants that sells any pizza that Ralph likes.

$$Q_2 := \pi_{\text{rname}}(\text{Sells} \bowtie \sigma_{\text{cname}=\text{'Ralph'}}(\text{Likes}))$$

This gives us the final answer of

$$Q_{\text{ans}} := Q_1 - Q_2$$

- (b) Let $\pi_A(R) = Q \cup Q'$ where $Q = R/S$ and Q' are the remaining tuples in $\pi_A()$ that are not in quotient of R/S .

The expression $\pi_A(R) \times S$ computes all the combinations of $\pi_A(R)$ and S .

Thus, $Q' = \pi_A((\pi_A(R) \times S) - R)$.

Therefore, $Q = \pi_A(R) - Q' = \pi_A(R) - \pi_A((\pi_A(R) \times S) - R)$.

2. Consider a relational database for a company that consists of the following two tables.

- **Offices**(office_id, building, level, room_number, area)
- **Employees**(emp_id, name, office_id, manager_id)

The database satisfies the following constraints:

1. **Offices** stores information about the office rooms in the company. Each room has a unique identifier **office_id** (*which is the primary key of **Offices***) and information on its building name, floor level, room number and floor area.
2. {building, level, room_number} is a *candidate key* of **Offices**.
3. **Employees** stores information about the employees in the company.
4. **emp_id** is the *primary key* of **Employees**.
5. The name of each employee must be a non-NULL value.
6. Each employee must be assigned to exactly one office identified by **office_id**.
7. Each employee may be managed by at most one manager.
8. If an employee is managed by someone, the **emp_id** of his/her manager is recorded in **manager_id**.
9. A record in **Offices** cannot be removed if there's some employee assigned to that office.
10. A manager in **Employees** cannot be removed if there's some other employee managed by that manager.
11. Any modification to **office_id** in **Offices** is propagated to other database records.
12. Any modification to **emp_id** in **Employees** is propagated to other database records.

Write SQL statements to create the database schema with appropriate attribute domains and constraints.

Solution:

```
1 DROP TABLE IF EXISTS Offices, Employees CASCADE;
2
3 CREATE TABLE Offices (
4     office_id    INTEGER,
5     building     TEXT NOT NULL,
6     level        INTEGER NOT NULL,
7     room_number  INTEGER NOT NULL,
8     area         INTEGER,
9     PRIMARY KEY (office_id),
10    UNIQUE (building, level, room_number)
11 );
12
13
```

```
14 CREATE TABLE Employees (  
15     emp_id      INTEGER,  
16     name        TEXT NOT NULL,  
17     office_id   INTEGER NOT NULL,  
18     manager_id  INTEGER,  
19     PRIMARY KEY (emp_id),  
20     FOREIGN KEY (office_id) REFERENCES Offices (office_id)  
21         ON UPDATE CASCADE,  
22     FOREIGN KEY (manager_id) REFERENCES Employees (emp_id)  
23         ON UPDATE CASCADE  
24 );
```

Note that constraints (6) and (7) are both enforced by the primary key constraint in `Employees` which ensure that there can't be two `Employees` records with the same primary key value and different values for `office_id` / `manager_id`.

3. Consider a relational database for a company that consists of the following two tables.

- **Books**(isbn, title, authors, year, edition, publisher, number_pages, price)
- **Customers**(cust_id, name, email)
- **Carts**(cust_id, isbn)
- **Purchase**(pid, purchase_date, cust_id)
- **Purchased_items**(pid, isbn)

The database satisfies the following constraints:

1. **Books** records information about the books available for sale in an online shop. Each book has a unique identifier `isbn` and information about its title, authors, publishers, publication year, edition, number of pages and selling price. The title and authors must have non-NULL values. The value of the edition must be non-NULL with one of the following values: *paperback*, *hardcover* or *ebook*. The selling price must have a positive value. If the number of pages is known, it must be a positive value.
2. **Customers** stores information about the shop's customers. Each customer has a unique identifier `cust_id`, a name and an email address. The name must have a non-NULL value.
3. **Carts** stores information about the books in customers' shopping carts. Each shopping cart record indicates a book that a customer is interested to purchase but has not yet purchased.
4. When a customer decides to purchase their selected books, a new record for this purchase is recorded in the **Purchase** table which has the following information:
 - (a) A unique identifier `pid`.
 - (b) The date of the purchase.
 - (c) The customer identifier.

In addition, each book in the customers' shopping cart is added to the **Purchased_items** table and the customers' shopping cart is emptied.

- (a) Write SQL statements to create a database schema with appropriate attribute domains and constraints.
- (b) Suppose that the schema of **Purchase** is changed with `purchase_date` being replaced by `purchase_timestamp`, where each customer has at most one purchase at any timestamp. How would this change affect your answer for part (a)?
- (c) For each of the following additional constraints on the database, state whether the constraint can be expressed using SQL constructs that you have learned. If it is possible, add this constraint to your database schema in part (a).
 - (i) If a book is a hardcover edition, its selling price must be at least 30.

- (ii) If a book has both hardcover and paperback editions (*for the same book title and authors*), the selling price for the hardcover edition must be higher than the selling price for the paperback edition.
 - (iii) If the number of pages in a book is more than 1000, the edition of the book must be an ebook or its price must be at least 100.
 - (iv) All the books published by 'Acme' from 2010 onwards have only ebook edition.
- (d) Consider the following additional constraints on the database:
- (i) If a customer is deleted from **Customers**, remove all the customers' records from **Carts** and **Purchase**.
 - (ii) If a book is deleted from **Books**, remove all records from **Carts** that reference this book. Additionally, for each of the records in **Purchased_items** that reference this book, change its **isbn** value to the default value of 0.
 - (iii) If a purchase is deleted from **Purchase**, remove all the records from **Purchased_items** that reference this purchase.
 - (iv) Any modification of **cust_id** in **Customers** is propagated to other records in the database.
 - (v) Any modification of **isbn** in **Books** is propagated to other records in the database.
 - (vi) Any modification of **pid** in **Purchase** is propagated to other records in the database.

Add these additional constraints to your answer for part (a)

Solution: Each table in a database schema must have a primary key.

(a) SQL statements

```

1 DROP TABLE IF EXISTS
2   Books, Customers, Carts, Purchase, Purchased_items
3   CASCADE;
4
5 CREATE TABLE Books (
6   isbn          TEXT,
7   title         TEXT NOT NULL,
8   authors       TEXT NOT NULL,
9   year          INTEGER,
10  edition       TEXT NOT NULL
11  CHECK (edition in ('hardcopy', 'paperback', 'ebook'))
12  ,
13  publisher      TEXT,
14  number_pages   INTEGER CHECK (number_pages > 0),
15  price          NUMERIC NOT NULL CHECK (price > 0),
16  PRIMARY KEY (isbn)
17 );
18

```

```

19
20 CREATE TABLE Customers (
21     cust_id    INTEGER,
22     name       TEXT NOT NULL,
23     email      TEXT NOT NULL,
24     PRIMARY KEY (cust_id)
25 );
26
27 CREATE TABLE Carts (
28     cust_id    INTEGER,
29     isbn       TEXT,
30     PRIMARY KEY (cust_id, isbn),
31     FOREIGN KEY (cust_id) REFERENCES Customers,
32     FOREIGN KEY (isbn) REFERENCES Books
33 );
34
35 CREATE TABLE Purchase (
36     pid        INTEGER,
37     purchase_date DATE NOT NULL,
38     cust_id    INTEGER NOT NULL,
39     PRIMARY KEY (pid),
40     FOREIGN KEY (cust_id) REFERENCES Customers
41 );
42
43 CREATE TABLE Purchased_Items (
44     pid    INTEGER,
45     isbn   TEXT,
46     PRIMARY KEY (pid, isbn),
47     FOREIGN KEY (pid) REFERENCES Purchase,
48     FOREIGN KEY (isbn) REFERENCES Books
49 );

```

(b) (purchase_timestamp, cust_id) is a candidate key of Purchase.

```

1 CREATE TABLE Purchase (
2     pid            INTEGER,
3     purchase_timestamp TIMESTAMP NOT NULL,
4     cust_id        INTEGER NOT NULL,
5     PRIMARY KEY (pid),
6     FOREIGN KEY (cust_id) REFERENCES Customers,
7     UNIQUE (purchase_timestamp, cust_id)
8 );

```

(c) The constraint of the form " $p \Rightarrow q$ " is equivalent to "(not p) or q ".

- (i) Can be expressed using a table constraint on Books:
CHECK ((edition <> 'hardcover') OR (price >= 30))
- (ii) Can **NOT** be expressed using the constructs learned so far.
- (iii) Can be expressed using a table constraint on Books:
CHECK ((number_pages <= 1000) OR (edition = 'ebook')
OR (price >= 100))
- (iv) Can be expressed using a table constraint on Books:
CHECK ((publisher <> 'Acme') OR (year < 2010)
OR (edition = 'ebook'))

(d) SQL statements

```

1  DROP TABLE IF EXISTS
2      Books, Customers, Carts, Purchase, Purchased_items
3      CASCADE;
4
5  CREATE TABLE Books (
6      isbn          TEXT,
7      title         TEXT NOT NULL,
8      authors       TEXT NOT NULL,
9      year          INTEGER,
10     edition       TEXT NOT NULL
11     CHECK (edition IN ('hardcopy', 'paperback', 'ebook'))
12     ,
13     publisher     TEXT,
14     number_pages  INTEGER CHECK (number_pages > 0),
15     price         NUMERIC NOT NULL CHECK (price > 0),
16     PRIMARY KEY (isbn)
17 );
18
19 CREATE TABLE Customers (
20     cust_id        INTEGER,
21     name          TEXT NOT NULL,
22     email         TEXT NOT NULL,
23     PRIMARY KEY (cust_id)
24 );
25
26 CREATE TABLE Carts (
27     cust_id        INTEGER,
28     isbn          TEXT,
29     PRIMARY KEY (cust_id, isbn),
30     FOREIGN KEY (cust_id) REFERENCES Customers
31     ON DELETE CASCADE ON UPDATE CASCADE,
32     FOREIGN KEY (isbn) REFERENCES Books
33     ON DELETE CASCADE ON UPDATE CASCADE
34 );
35
36 CREATE TABLE Purchase (
37     pid           INTEGER,
38     purchase_date DATE NOT NULL,
39     cust_id       INTEGER NOT NULL,
40     PRIMARY KEY (pid),
41     FOREIGN KEY (cust_id) REFERENCES Customers
42     ON DELETE CASCADE ON UPDATE CASCADE
43 );
44
45 CREATE TABLE Purchased_Items (
46     pid          INTEGER,
47     isbn         TEXT DEFAULT '0',
48     PRIMARY KEY (pid, isbn),
49     FOREIGN KEY (pid) REFERENCES Purchase
50     ON DELETE CASCADE ON UPDATE CASCADE,
51     FOREIGN KEY (isbn) REFERENCES Books
52     ON DELETE SET DEFAULT ON UPDATE CASCADE
53 );

```

For the "ON DELETE SET DEFAULT" action to work in Purchased_items

when a referenced book in **Books** is deleted, there must exist a record in **Books** with `isbn = '0'`. Otherwise, the deletion operation will be rejected.