

1. Consider a database consisting of the following two tables shown below:

bar		foo	
a	b	f	a
1	10	100	2
2	20	200	7
3	30	300	3
4	40	400	2

For each of the following queries on the database, either state that the query is an *invalid* SQL query or show the query's output if the query is a *valid* SQL query.

(a) Query A

```
1 SELECT *
2 FROM   bar b
3 WHERE EXISTS (
4     SELECT 1
5     FROM   foo f
6     WHERE  f.f > 100
7     AND    f.a = b.a
8 );
```

(c) Query C

```
1 SELECT *
2 FROM   bar b
3 WHERE EXISTS (
4     SELECT 1
5     FROM   foo f
6     WHERE  f.f > 100
7     AND    a = b.a
8 );
```

(b) Query B

```
1 SELECT *
2 FROM   bar b
3 WHERE EXISTS (
4     SELECT 1
5     FROM   foo f
6     WHERE  f.f > 100
7 )
8 AND    f.a = b.a;
```

(d) Query D

```
1 SELECT *
2 FROM   bar b
3 WHERE EXISTS (
4     SELECT 1
5     FROM   foo f
6     WHERE  f.f > 100
7     AND    a = a
8 );
```

(e) Query E

```
1 SELECT *
2 FROM   bar b
3 WHERE EXISTS (
4     SELECT 1
5     FROM   foo f
6     WHERE  f.f > 100
7     AND    f.a = b.a
8     AND    b > 20
9 );
```

**Solution:** We begin the solution with a warning:

**⚠** *It is a good defensive programming practice to use distinct aliases and use explicitly qualified column names.*

- (a) **Valid:** This follows the good practice.

a	b
2	20
3	30

- (b) **Invalid:** The alias **f** for table **foo** declared in the *inner query* is not visible to the *outer query* due to scoping.

- (c) **Valid:** The expression "**a = b.a**" is *equivalent* to "**f.a = b.a**".

a	b
2	20
3	30


- (d) **Valid:** The expression "**f > 100**" is *equivalent* to "**f.f > 100**" (*the f refers to the attribute not the table alias even though both have the same name*) and "**a = a**" is *equivalent* to "**f.a = f.a**" (*which is always true, hence the result*).

a	b
1	10
2	20
3	30
4	40

- (e) **Valid:** The expression "**b > 20**" is *equivalent* to "**b.b > 20**" (*the b refers to the attribute not the table alias even though both have the same name*). How does it know that it is "**b.b**" instead of "**f.b**"? SQL is smart that way, it knows that the table **foo** or **f** does not have an attribute **b**.

a	b
3	30

2. Questions 2 to 5 are based on the pizza database schema used in the lectures; we show its ER diagram below.



For each of the following queries, write an equivalent SQL query that does not use any subquery.

(a) Query A

```

1 SELECT DISTINCT cname
2 FROM   Likes L
3 WHERE EXISTS (
4     SELECT 1
5     FROM   Sells S
6     WHERE S.rname = 'Corleone Corner'
7         AND S.pizza = L.pizza
8 );
  
```

(b) Query B

```

1 SELECT cname
2 FROM   Customers C
3 WHERE NOT EXISTS (
4     SELECT 1
5     FROM   Likes L, Sells S
6     WHERE S.rname = 'Corleone Corner'
7         AND S.pizza = L.pizza
8         AND C.cname = L.cname
9 );
  
```

(c) Query C

```

1 SELECT DISTINCT rname
2 FROM   Sells
3 WHERE rname <> 'Corleone Corner'
4     AND price > ANY (
5     SELECT price
6     FROM   Sells
7     WHERE rname = 'Corleone Corner'
8 );
  
```

(d) Query D

```

1 SELECT rname, pizza, price
2 FROM   Sells S
3 WHERE price >= ALL (
4     SELECT S2.price
5     FROM   Sells S2
6     WHERE S2.rname = S.rname
7         AND S2.price IS NOT NULL
8 );
  
```

**Solution:**

## (a) Query A

```
1 SELECT DISTINCT cname
2 FROM Likes L, Sells S
3 WHERE S.rname = 'Corleone Corner'
4 AND S.pizza = L.pizza;
```

## (b) Query B

```
1 SELECT cname FROM Customers
2 EXCEPT
3 SELECT cname
4 FROM Likes L, Sells S
5 WHERE S.rname = 'Corleone Corner'
6 AND S.pizza = L.pizza;
```


## (c) Query C

```
1 SELECT DISTINCT S.rname
2 FROM Sells S, Sells S2
3 WHERE S.rname <> 'Corleone Corner'
4 AND S2.rname = 'Corleone Corner'
5 AND S.price > S2.price;
```

## (d) Query D

```
1 SELECT rname, pizza, price
2 FROM Sells
3 WHERE price IS NOT NULL
4 EXCEPT
5 SELECT S.rname, S.pizza, S.price
6 FROM Sells S, Sells S2
7 WHERE S.rname = S2.rname
8 AND S.price < S2.price;
```

3. Write an SQL query to answer each of the following questions on the pizza database *without using aggregate functions*. Remove duplicate records from all query results.



- Find pizzas that Alice likes but Bob does not like.
- Find pizzas that are sold by at most one restaurant in each area; exclude pizzas that are not sold by any restaurant.
- Find all tuples  $(A, P, P_{min})$  where  $P$  is a pizza that is available in area  $A$  (i.e., there is some restaurant in area  $A$  selling pizza  $P$ ) and  $P_{min}$  is the lowest price of  $P$  in area  $A$ .
- Find all tuples  $(A, P, P_{min}, P_{max})$  where  $P$  is a pizza that is available in area  $A$  (i.e., there is some restaurant in area  $A$  selling pizza  $P$ ),  $P_{min}$  is the lowest price of  $P$  in area  $A$  and  $P_{max}$  is the highest price of  $P$  in area  $A$ .

### Solution:

#### (a) Solution 1:

```

1 SELECT pizza FROM Likes
2 WHERE cname = 'Alice'
3    AND pizza NOT IN (
4        SELECT pizza FROM Likes
5        WHERE cname = 'Bob'
6    );
  
```

#### Solution 2:

```

1 SELECT pizza FROM Likes L1
2 WHERE cname = 'Alice'
3    AND NOT EXISTS (
4        SELECT 1 FROM Likes L2
5        WHERE L2.cname = 'Bob' AND L2.pizza = L1.pizza
6    );
  
```

#### Solution 3:

```

1 SELECT pizza FROM Likes
2 WHERE cname = 'Alice'
3    AND NOT pizza = ANY (
4        SELECT pizza FROM Likes
5        WHERE cname = 'Bob'
6    );
  
```

Note that if Bob does not like any pizza, then the ANY subquery will evaluate to FALSE and NOT FALSE will evaluate to TRUE. Thus the query will return all the pizzas that Alice likes.

**Solution 4:** *Probably simplest*

```

1 SELECT pizza FROM Likes WHERE cname = 'Alice'
2 EXCEPT
3 SELECT pizza FROM Likes WHERE cname = 'Bob';

```

**WRONG ANSWER:** *The following answer is incorrect*

```

1 SELECT pizza FROM Likes
2 WHERE cname = 'Alice'
3 AND pizza <> ANY (
4     SELECT pizza FROM Likes
5     WHERE cname = 'Bob'
6 );

```

This answer looks similar to **Solution 3** BUT it is *incorrect*. If Bob does not like any pizza, then the ANY subquery will evaluate to FALSE and the query will return an *empty set*. This will be incorrect if Alice likes some pizza.

- (b) A pizza is the output if there does not exist two distinct restaurants that are located in the same area selling that pizza.

```

1 SELECT DISTINCT pizza
2 FROM Sells S3
3 WHERE NOT EXISTS (
4     SELECT 1
5     FROM Sells S, Restaurants R, Sells S2, Restaurants
6         R2
7     WHERE S.rname = R.rname AND S2.rname = R2.rname
8           AND S.pizza = S2.pizza AND R.area = R2.area
9           AND R.rname <> R2.rname AND S.pizza = S3.pizza
10 );

```

**WRONG ANSWER:** *The following answer is incorrect*

```

1 SELECT DISTINCT pizza
2 FROM Sells S, Restaurants R
3 WHERE S.rname = R.rname
4 AND NOT EXISTS (
5     SELECT 1
6     FROM Sells S2, Restaurants R2
7     WHERE S2.rname = R2.rname AND S.pizza = S2.pizza
8           AND R.area = R2.area AND R.rname <> R2.rname
9 );

```

This answer computes the pizzas that are sold by at most one restaurant in *some* area, which is a weaker condition than what is required by the question.

(c) A possible solution

```

1  SELECT DISTINCT R.area, S.pizza, S.price
2  FROM    Restaurants R, Sells S
3  WHERE   R.rname = S.rname
4         AND S.price <= ALL (
5         SELECT S2.price
6         FROM    Restaurants R2, Sells S2
7         WHERE   R2.rname = S2.rname
8             AND R2.area = R.area
9             AND S2.pizza = S.pizza
10 );

```

(d) You should recognize that this query is simply an extension of the previous query requiring an additional information (*i.e.*, highest selling price) for each area-pizza pair. For a given area-pizza pair  $(A, P)$ , the following query will compute the highest price of pizza  $P$  in area  $A$ .

```

1  SELECT DISTINCT S2.price
2  FROM    Restaurants R2, Sells S2
3  WHERE   R2.rname = S2.rname
4         AND R2.area = A AND S2.pizza = P
5         AND R2.price <= ALL (
6         SELECT S2.price
7         FROM    Restaurants R3, Sells S3
8         WHERE   R2.rname = S2.rname
9             AND R3.area = A
10             AND S3.pizza = P
11 );

```

Since the above query will return a *single one-column tuple*, it can be used as a **scalar** subquery to extend the previous question's solution as follows:

```

1  SELECT DISTINCT R.area, S.pizza, S.price AS minPrice, (
2      SELECT DISTINCT S2.price
3      FROM    Restaurants R2, Sells S2
4      WHERE   R2.rname = S2.rname
5          AND R2.area = A AND S2.pizza = P
6          AND R2.price <= ALL (
7              SELECT S2.price
8              FROM    Restaurants R3, Sells S3
9              WHERE   R2.rname = S2.rname
10                 AND R3.area = A
11                 AND S3.pizza = P
12          ) AS maxPrice
13  FROM    Restaurants R, Sells S
14  WHERE   R.rname = S.rname
15         AND S.price <= ALL (
16         SELECT S2.price
17         FROM    Restaurants R2, Sells S2
18         WHERE   R2.rname = S2.rname
19             AND R2.area = R.area
20             AND S2.pizza = S.pizza
21 );

```

We will learn about other (*simpler and more elegant*) solutions for such queries later in class.



4. Consider the query to find distinct restaurants that are located in the East area. The following are two possible SQL answers (*denoted by  $Q_1$  and  $Q_2$* ) for this query.


$Q_1$  Query 1

```
1 SELECT DISTINCT S.rname
2 FROM Sells S JOIN Restaurants R
3 ON S.rname = R.rname AND R.area = 'East';
```

$Q_2$  Query 2

```
1 SELECT DISTINCT S.rname
2 FROM Sells S, Restaurants R
3 WHERE S.rname = R.rname
4 AND R.area = 'East';
```

The semantics of these two SQL queries are defined by the relational algebra expressions shown below. Discuss whether  $Q_1$  and  $Q_2$  are equivalent queries.



**Solution:** Queries  $Q_1$  and  $Q_2$  are **equivalent**. Whether the selection predicate " $area = 'East'$ " is evaluated *before* or *after* the join/cross product operation does not change the semantics of the query.

5. Consider the query to find distinct restaurants that are located in the East area or restaurants that sell some pizza that Lisa likes, where the restaurants that do not sell any pizza are to be excluded. The following are two possible SQL answers (*denoted by  $Q_1$  and  $Q_2$* ) for this query.

$Q_1$  Query 1

```

1 SELECT DISTINCT S.rname
2 FROM Sells S JOIN Restaurants R
3 ON S.rname = R.rname AND R.area = 'East'
4 UNION
5 SELECT DISTINCT S.rname
6 FROM Sells S JOIN Likes L
7 ON S.pizza = L.pizza AND L.cname = 'Lisa';

```


$Q_2$  Query 2

```

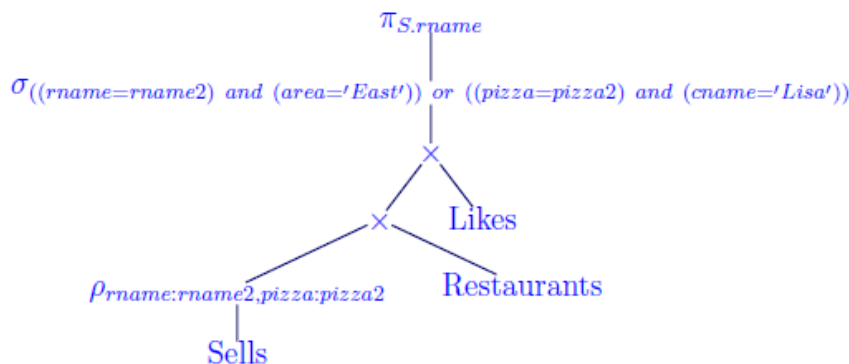
1 SELECT DISTINCT S.rname
2 FROM Sells S, Restaurants R, Likes L
3 WHERE (S.rname = R.rname AND R.area = 'East')
4 OR (S.pizza = L.pizza AND L.cname = 'Lisa');

```

The semantics of these two SQL queries are defined by the relational algebra expressions shown below. Discuss whether  $Q_1$  and  $Q_2$  are equivalent queries.



(a)  $Q_1$



(b)  $Q_2$

**Solution:** Queries  $Q_1$  and  $Q_2$  are **NOT equivalent**. Observe that if the Likes relation is *empty*, then query  $Q_1$  simplifies to

$$\pi_{\text{rname}}(\text{Sells} \bowtie_{\sigma_{\text{area}='East'}}(\text{Restaurants}))$$

BUT the result of  $Q_2$  is always an empty set due to

$$\text{Sells} \times \text{Restaurants} \times \emptyset$$

6. Consider *again* the following relational schema discussed in Tutorial 2.

```
1 CREATE TABLE Offices (  
2     office_id    INTEGER,  
3     building     TEXT NOT NULL,  
4     level        INTEGER NOT NULL,  
5     room_number  INTEGER NOT NULL,  
6     area         INTEGER,  
7     PRIMARY KEY (office_id),  
8     UNIQUE (building, level, room_number)  
9 );  
10  
11 CREATE TABLE Employees (  
12     emp_id       INTEGER,  
13     name         TEXT NOT NULL,  
14     office_id    INTEGER NOT NULL,  
15     manager_id   INTEGER,  
16     PRIMARY KEY (emp_id),  
17     FOREIGN KEY (office_id) REFERENCES Offices (office_id)  
18         ON UPDATE CASCADE,  
19     FOREIGN KEY (manager_id) REFERENCES Employees (emp_id)  
20         ON UPDATE CASCADE  
21 );
```

Suppose that the office with `office_id = 123` needs to be renovated. Write an SQL statement to reassign the employees located in this office to another temporary office located at room number 11 on level 5 at the building named *Tower1*.

**Solution:** SQL update statement.

```
1 UPDATE Employees  
2 SET     office_id = (SELECT office_id FROM Offices  
3                     WHERE building = 'Tower1'  
4                     AND level = 5  
5                     AND room_number = 11)  
6 WHERE  office_id = 123;
```

7. Given the tables  $R$  and  $S$  shown below, compute the output of each of the following queries.

- (a) `SELECT * FROM R NATURAL JOIN S;`
- (b) `SELECT * FROM R INNER JOIN S ON R.A = S.A;`
- (c) `SELECT * FROM R LEFT OUTER JOIN S ON R.A = S.A;`
- (d) `SELECT * FROM R RIGHT OUTER JOIN S ON R.A = S.A;`
- (e) `SELECT * FROM R FULL OUTER JOIN S ON R.A = S.A;`

R					S			
X	A	Y	B	Z	A	B	C	D
0	10	0	9	2	17	1	20	100
30	8	0	5	1	4	2	40	200
60	4	1	3	3	4	3	30	100
0	0	0	4	5	8	5	60	500

**Solution:**

- (a) `SELECT * FROM R NATURAL JOIN S;`

A	B	X	Y	Z	C	D
8	5	30	0	1	60	100
4	3	60	1	3	30	100

- (b) `SELECT * FROM R INNER JOIN S ON R.A = S.A;`

X	A	Y	B	Z	A	B	C	D
30	8	0	5	1	8	5	60	500
60	4	1	3	3	4	2	40	200
60	4	1	3	3	4	3	30	100

- (c) `SELECT * FROM R LEFT OUTER JOIN S ON R.A = S.A;`

X	A	Y	B	Z	A	B	C	D
0	10	0	9	2	NULL	NULL	NULL	NULL
30	8	0	5	1	8	5	60	500
60	4	1	3	3	4	2	40	200
60	4	1	3	3	4	3	30	100
90	0	0	4	5	NULL	NULL	NULL	NULL

- (d) `SELECT * FROM R RIGHT OUTER JOIN S ON R.A = S.A;`

X	A	Y	B	Z	A	B	C	D
30	8	0	5	1	8	5	60	500
60	4	1	3	3	4	2	40	200
60	4	1	3	3	4	3	30	100
NULL	NULL	NULL	NULL	NULL	17	1	20	100

(e) `SELECT * FROM R FULL OUTER JOIN S ON R.A = S.A;`

X	A	Y	B	Z	A	B	C	D
0	10	0	9	2	NULL	NULL	NULL	NULL
30	8	0	5	1	8	5	60	500
60	4	1	3	3	4	2	40	200
60	4	1	3	3	4	3	30	100
90	0	0	4	5	NULL	NULL	NULL	NULL
NULL	NULL	NULL	NULL	NULL	17	1	20	100