

1. Consider the following relational schema consisting of two tables.

```
1 CREATE TABLE R (  
2   a   INTEGER PRIMARY KEY,  
3   b   INTEGER,  
4   c   INTEGER  
5 );  
6  
7 CREATE TABLE S (  
8   x   INTEGER PRIMARY KEY,  
9   y   INTEGER REFERENCES R(a)  
10 );
```

For each of the following queries on the database, state whether the query is a valid SQL query.

- (a) Query A

```
1 SELECT  a, b, SUM(c)  
2 FROM    R JOIN S ON R.a = S.y  
3 GROUP BY a, b;
```

- (b) Query B

```
1 SELECT  b, SUM(c)  
2 FROM    R JOIN S ON R.a = S.y  
3 GROUP BY a, b;
```

- (c) Query C

```
1 SELECT  a, b, SUM(c)  
2 FROM    R JOIN S ON R.a = S.y  
3 GROUP BY a;
```

- (d) Query D

```
1 SELECT  a, b, x, SUM(c)  
2 FROM    R JOIN S ON R.a = S.y  
3 GROUP BY a, b;
```

- (e) Query E

```
1 SELECT  a, b, y, SUM(c)  
2 FROM    R JOIN S ON R.a = S.y  
3 GROUP BY a, b;
```

- (f) Query F

```
1 SELECT  a, b  
2 FROM    R JOIN S ON R.a = S.y  
3 GROUP BY a, b  
4 HAVING  SUM(c) > 10;
```

- (g) Query G

```
1 SELECT  a, b  
2 FROM    R JOIN S ON R.a = S.y  
3 GROUP BY a, b  
4 HAVING  SUM(x) > 10;
```

(h) Query H

```

1 SELECT    a, b
2 FROM      R JOIN S ON R.a = S.y
3 WHERE     SUM(c) > 10;

```

(i) Query I

```

1 SELECT    a, b
2 FROM      R JOIN S ON R.a = S.y
3 GROUP BY  a, b
4 WHERE     SUM(x) > SUM(c);

```

(j) Query J

```

1 SELECT    b, SUM(c)
2 FROM      R;

```

**Solution:**

- (a) **Valid.**
- (b) **Valid:** It is fine for an attribute in the **GROUP BY** clause to be absent from the **SELECT** clause.
- (c) **Valid:** It is fine for attribute **b** to be in the **SELECT** clause since the *primary key* of **R** is included in the **GROUP BY** clause.
- (d) **Invalid:** Attribute **x** is missing from the **GROUP BY** clause.
- (e) **Invalid:** Attribute **y** and the primary key of **S** are missing from the **GROUP BY** clause.
- (f) **Valid.**
- (g) **Valid.**
- (h) **Invalid:** Aggregate function cannot be used in the **WHERE** clause.
- (i) **Invalid:** **HAVING** should be used instead.
- (j) **Invalid:** A query that has no **GROUP BY** clause cannot have both aggregate function and non-aggregate attribute in the **SELECT** clause.

2. Consider the same relational schema from the previous question.

```
1 CREATE TABLE R (  
2   a   INTEGER PRIMARY KEY,  
3   b   INTEGER,  
4   c   INTEGER  
5 );  
6  
7 CREATE TABLE S (  
8   x   INTEGER PRIMARY KEY,  
9   y   INTEGER REFERENCES R(a)  
10 );
```

Are these two queries *equivalent*?

(a)  $Q_1$

```
1 SELECT   COUNT(c)  
2 FROM     R  
3 WHERE    a = 10;
```


(b)  $Q_2$

```
1 SELECT   COUNT(c)  
2 FROM     R  
3 WHERE    a = 10  
4 GROUP BY a;
```

**Solution:**  $Q_1$  and  $Q_2$  are **NOT** equivalent queries. Consider the case where there is no record in R with  $a = 10$ .

- (a)  $Q_1$ : The WHERE clause will evaluate to an *empty table* and the aggregate function COUNT(c) will evaluate to a *single-row table* with a *single-column* value of 0.
- (b)  $Q_2$ : The WHERE clause will evaluate to an *empty table* so the GROUP BY clause will **NOT** produce any group. As such, the COUNT(c) will **NOT** be evaluated at all. As a result, the query result will be an *empty table*.

3. This question is based on the pizza database schema shown below.



Answer each of the following queries using SQL. For parts (a) to (e), remove duplicate records from all query results.

- Find the most expensive pizzas and the restaurants that sell them (*at the most expensive price*). Do **NOT** use any aggregate function in your answer.
- Find all restaurant pairs  $(R_1, R_2)$  such that the price of the most expensive pizza sold by  $R_1$  is *higher* than that of  $R_2$ . Exclude restaurant pairs where  $R_2$  do not sell any pizza.
- For each restaurant that sells some pizza, find the restaurant name and the average price of its pizzas if its average price is higher than \$22. Do **NOT** use the HAVING clause in your answer.
- For each restaurant  $R$  that sells some pizza, let  $totalPrice(R)$  denote the total price of all the pizzas sold by  $R$ . Find all pairs  $(R, totalPrice(R))$  where  $totalPrice(R)$  is *higher* than the average of  $totalPrice()$  over all the restaurants.
- Find the customer pairs  $(C_1, C_2)$  such that  $C_1 < C_2$  and they like *exactly* the same pizzas. Exclude customer pairs that do not like any pizza. Do **NOT** use the EXCEPT operator in your answer.
- Write an SQL statement to *increase* the selling prices of pizzas as follows:
  - Increase by \$3 if the restaurant is located in 'Central'.
  - Increase by \$2 if the restaurant is located in 'East'.
  - Increase by \$1 *otherwise*.

### Solution:

#### (a) Solution 1:

```

1 SELECT pizza, rname
2 FROM   Sells
3 WHERE  price >= ALL (SELECT price FROM Sells);
  
```

#### Solution 2:

```

1 SELECT pizza, rname
2 FROM   Sells S1
3 WHERE  NOT EXISTS (
4     SELECT 1
5     FROM   Sells S2
6     WHERE  S2.price > S1.price
7 );
  
```

(b) Solution 1:

```

1 SELECT R1.rname, R2.rname
2 FROM   Restaurants R1, Restaurants R2
3 WHERE  (SELECT MAX(price) FROM Sells
4         WHERE rname = R1.rname)
5         >
6         (SELECT MAX(price) FROM Sells
7         WHERE rname = R2.rname);

```

Solution 2:

```

1 WITH RestMaxPrice AS (
2     SELECT rname, (SELECT MAX(price) FROM Sells
3                   WHERE rname = R.rname) as maxPrice
4     FROM   Restaurants R
5 )
6 SELECT R1.rname, R2.rname
7 FROM   RestMaxPrice R1, RestMaxPrice R2
8 WHERE  R1.maxPrice > R2.maxPrice;

```

## (c) A possible solution

```

1 WITH RestAvgPrice AS (
2     SELECT  rname, AVG(price) as avgPrice
3     FROM    Sells
4     GROUP BY rname
5 )
6 SELECT *
7 FROM   RestAvgPrice
8 WHERE  avgPrice > 22;

```

(d) Solution 1:

```

1 WITH RestTotalPrice AS (
2     SELECT  rname, SUM(price) as totalPrice
3     FROM    Sells
4     GROUP BY rname
5 )
6 SELECT rname, totalPrice
7 FROM   RestTotalPrice
8 WHERE  totalPrice > (SELECT AVG(totalPrice) FROM
9                     RestTotalPrice);

```

Solution 2:

```

1 SELECT  rname, SUM(price) as totalPrice
2 FROM    Sells S
3 GROUP BY rname
4 HAVING  SUM(price) > (SELECT SUM(price) / COUNT(
5                       DISTINCT rname) FROM Sells);

```

WRONG ANSWER 1: The following query is an **invalid** query

```

1 SELECT  rname, SUM(price) as totalPrice
2 FROM    Sells S
3 GROUP BY rname
4 HAVING  totalPrice > (SELECT SUM(price) / COUNT(
5                       DISTINCT rname) FROM Sells);

```

`totalPrice` is *undefined* in the HAVING clause as the SELECT clause is conceptually evaluated *after* the HAVING clause.

**WRONG ANSWER 2:** The following query produce the wrong result

```
1 SELECT rname, SUM(price)
2 FROM Sells
3 GROUP BY rname
4 HAVING SUM(price) > SUM(price) / COUNT(*);
```

The query above is *incorrect* because both *SUM(price)* and *COUNT(\*)* in the HAVING clause are computed w.r.t. a group.

- (e) **Solution 1:** Customer C1 and C2 like exactly the same pizzas if and only if (a) for every pizza that C1 likes, C2 also likes that pizza, and (b) for every pizza that C2 likes, C1 also likes that pizza. Conditions (a) holds if there does not exists any pizza that C1 likes that C2 does not like.

```
1 SELECT C1.cname, C2.cname
2 FROM Customers C1, Customers C2
3 WHERE C1.cname < C2.cname
4 AND EXISTS (SELECT 1 FROM Likes L WHERE L.cname = C1.
5             cname)
5 AND NOT EXISTS (
6     SELECT 1
7     FROM Likes L1
8     WHERE L1.cname = C2.cname
9           AND NOT EXISTS (
10                SELECT 1
11                FROM Likes L2
12                WHERE L2.cname = C2.cname
13                      AND L2.pizza = L1.pizza
14            )
15 )
16 AND NOT EXISTS (
17     SELECT 1
18     FROM Likes L2
19     WHERE L2.cname = C2.cname
20           AND NOT EXISTS (
21                SELECT 1
22                FROM Likes L1
23                WHERE L1.cname = C1.cname
24                      AND L1.pizza = L2.pizza
25            )
26 );
```

**Solution 2:** This solution is based on the property that if customer C1 likes the set of pizzas  $S1$  and customer C2 likes the set of pizzas  $S2$ , then  $S1 = S2$  iff  $|S1 \cap S2| = |S1| = |S2|$  where  $|X|$  denote the cardinality of set  $X$ . Alternatively, you can also use the following equality  $|S1 \cap S2| = |S1 \cup S2|$ .

Here, we use CTE to solve the problem but a solution without CTE is possible.

```

1  WITH NumLike AS (
2      SELECT  cname, COUNT(*) AS num FROM Likes L
3      GROUP BY cname
4  ), NumBothLike AS (
5      SELECT  L1.cname AS cname1, L2.cname AS cname2, COUNT
6              (*) AS num
7      FROM    Likes L1, Likes L2
8      WHERE   (L1.pizza = L2.pizza) AND (L1.cname < L2.
9              cname)
10     GROUP BY L1.cname, L2.cname
11 )
12 SELECT  cname1, cname2
13 FROM    NumBothLike B
14 WHERE   num = (SELECT num FROM NumLike N WHERE N.cname =
15              B.cname1)
16         AND num = (SELECT num FROM NumLike N WHERE N.cname =
17              B.cname2);

```

**Solution 3:** This solution is based on the property that if customers  $C1$  and  $C2$  like exactly the same pizzas, then for the subset of all tuples  $S \subseteq \text{Likes}$  that are associated with  $C1$  or  $C2$ , there must be exactly two tuples in  $S$  associated with each distinct pizza in  $S$ .

```

1  SELECT  C1.cname, C2.cname
2  FROM    Customers C1, Customers C2
3  WHERE   C1.cname < C2.cname
4         AND C1.cname IN (SELECT cname FROM Likes)
5         AND NOT EXISTS (
6             SELECT  1
7             FROM    Likes
8             WHERE   cname IN (C1.cname, C2.cname)
9             GROUP BY pizza
10             HAVING  COUNT(*) <> 2
11         );

```

(f) **Solution:** This is simply done using the UPDATE statement with CASE expression.

```

1  UPDATE Sells S
2  SET   price = CASE (
3          SELECT  area
4          FROM    Restaurants
5          WHERE   rname = S.rname
6      )
7      WHEN 'Central' THEN price + 3
8      WHEN 'East' THEN price + 2
9      ELSE price + 1
10 END;

```

4. In the lecture, we have discussed an algebraic approach to derive a SQL query involving universal quantification. In this question, we consider another approach to derive the same SQL query based on predicate logic.

The following predicate logic expression specifies the set of names of students who have enrolled in every course offered by the CS department.

$$\{S.name \mid S \in Students \wedge \forall C(C \in Courses \wedge C.dept = 'CS' \implies \exists E(E \in Enrolls \wedge E.sid = S.studentId \wedge E.cid = C.courseId))\}$$

Rewrite the above expression into an equivalent expression without using any universal quantifier and implication operators, and translate it into an equivalent SQL query.

**Solution:** The query is read as

A set of students who have taken **ALL** courses (*i.e.*, modules) offered by CS department.

The **ALL** in bold is the universal quantification (*i.e.*, called *forall* with the symbol  $\forall$ ). We can remove this by rewriting the query as:

A set of students such that there **IS NO** courses (*i.e.*, modules) offered by CS department that the student has not taken.

We have now remove universal quantification and replace it with existential quantification but negated. The new query now can be written as follows:

```

1 SELECT S.name
2 FROM Students S
3 WHERE NOT EXISTS (
4     SELECT 1
5     FROM Courses C
6     WHERE C.dept = 'CS'
7     AND NOT EXISTS (
8         SELECT 1
9         FROM Enrolls E
10        WHERE E.sid = S.studentId
11              AND E.cid = C.courseId
12     )
13 );
```

Another way to read this is to see it logically with set comprehension using double negation:

1.  $\forall C(C \in Courses \wedge C.dept = 'CS' \implies \exists E(E \in Enrolls \wedge E.sid = S.studentId \wedge E.cid = C.courseId))$
2.  $\neg\neg\forall C(C \in Courses \wedge C.dept = 'CS' \implies \exists E(E \in Enrolls \wedge E.sid = S.studentId \wedge E.cid = C.courseId))$  (double negation)



3.  $\neg\exists C\neg(C \in Courses \wedge C.dept = 'CS' \implies \exists E(E \in Enrolls \wedge E.sid = S.studentId \wedge E.cid = C.courseId))$   
(De Morgan's law)
4.  $\neg\exists C\neg(\neg(C \in Courses \wedge C.dept = 'CS') \vee (\exists E(E \in Enrolls \wedge E.sid = S.studentId \wedge E.cid = C.courseId)))$   
(definition of  $\implies$ )
5.  $\neg\exists C((C \in Courses \wedge C.dept = 'CS') \wedge (\neg\exists E(E \in Enrolls \wedge E.sid = S.studentId \wedge E.cid = C.courseId)))$   
(De Morgan's law)

From here, we can implement  $\neg\exists$  can be implemented as NOT EXISTS in SQL.