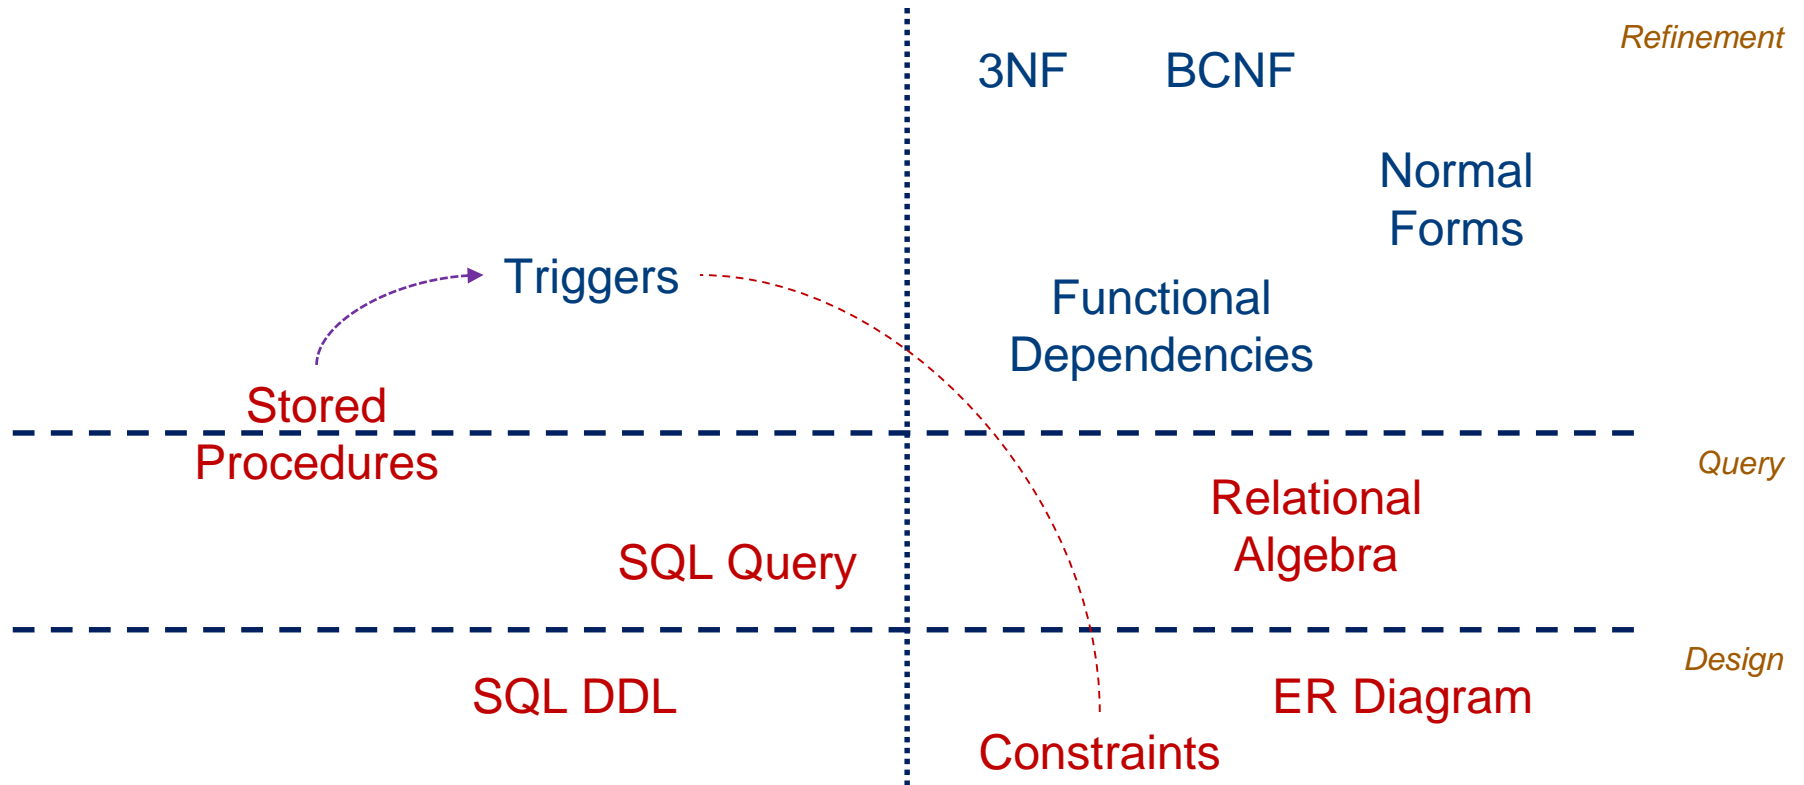


CS2102 Database Systems

Lecture 9 – Triggers

Roadmap



Previous Lectures

- **SQL Functions and Procedures**

- Named operations
- Expressive control structures (*Turing complete*)

- **Constraints**

- Key constraints
- Referential constraints
- Uniqueness
- Check

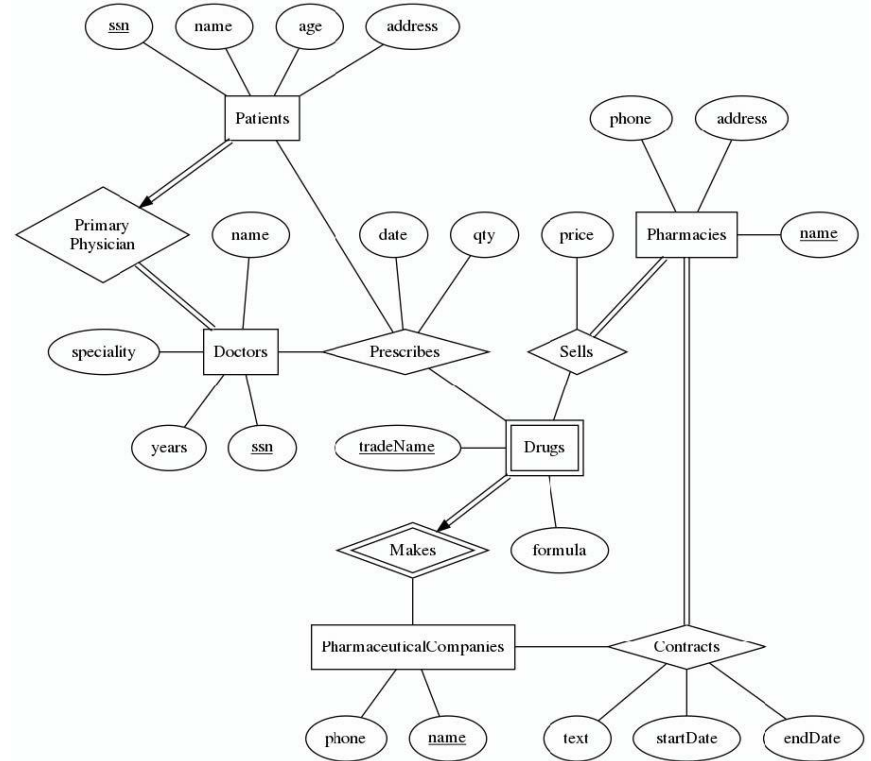
- What if there are constraints that cannot be handled by the constructs above?

Difficult Constraints

- **Tutorial 03**

- There is exactly one contract between a pharmacy and a pharmaceutical company if and only if that pharmacy sells some drug that is made by that pharmaceutical company.

This is quite complex, instead, we will use a simpler motivating example but Trigger is powerful enough to handle this




Difficult Constraints

- **ISA Hierarchy**

- A different difficult constraint that cannot be enforced via relational schema only is the ISA hierarchy with non-overlapping constraint (i.e., does not satisfy overlapping constraint).

The solution is:

- For each insertion to E1
- Check if the row is already present in E2
- Same for insert to E2, check on E1
- Checking can be done via trigger



Triggers

Triggers

- Other Motivating Example

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered
 - This should be done automatically
 - The user of the database should not be bothered to write SQL statement to insert
 - This should be done *each time* an insertion occurs regardless of how it is done
 - We want to record the following data
 - The name of the student
 - The date of entry

Triggers

- **Other Motivating Example**

What we want?

1. Automatic insertion
2. Regardless of how

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered

- **Idea?**

- A procedure that enters both data
 - ❖ What if the user forgot to use the procedure?

```
CREATE OR REPLACE PROCEDURE enter_data
    (Name TEXT, Mark INT)
AS $$

    INSERT INTO Scores VALUES (Name, Mark);
    INSERT INTO Scores_Log
        VALUES (Name, CURRENT_DATE);

$$ LANGUAGE sql;
```


Triggers

- Other Motivating Example

Issues?

1. How to check insertion?
2. How to get the name?

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47




Table "Scores_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered

- What we want?

- A procedure to insert into Scores_Log
- But called *automatically*

```
CREATE OR REPLACE PROCEDURE log_score()  
AS $$  
BEGIN  
    IF (there is an insertion into Scores) THEN  
        INSERT INTO Scores_Log  
            VALUES (Name, CURRENT_DATE);  
    END IF;  
END;  
$$ LANGUAGE plpgsql;
```

① How to check insertion?

② How to get the name?

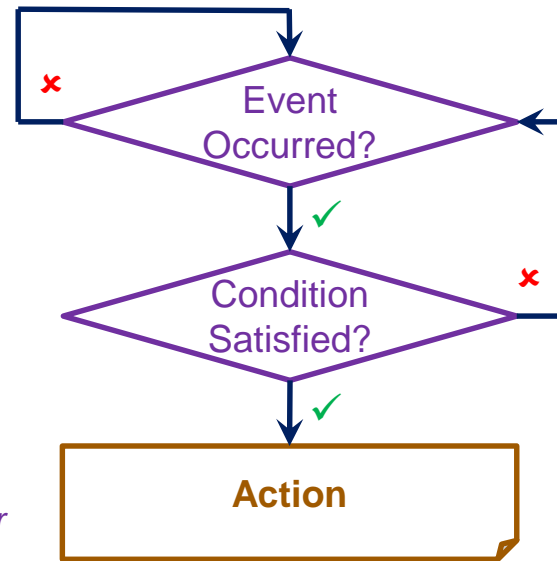
Triggers

- **Basic of Triggers**

- Trigger is an **event-condition-action** (ECA) rule
 - When **event** occurs
 - Test **condition**
 - If satisfied, execute **action**

- **Example**

- **Event** New tuple inserted into Scores
 - **Condition** Nothing (*always execute action*)
 - **Action** Insert into Scores_Log
- } *Trigger*
- } *Trigger Function*



Triggers

- Other Motivating Example

Trigger

```
CREATE TRIGGER score_log  
AFTER INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION  
log_score();
```

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

- Suppose we want to log when the marks are entered

- Trigger Function

- Actions to run
when event
occurred and
conditions
satisfied

```
CREATE OR REPLACE FUNCTION log_score()  
RETURNS TRIGGER AS $$  
BEGIN  
  
    INSERT INTO Scores_Log  
        VALUES (NEW.Name, CURRENT_DATE);  
    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

Triggers

- The Trigger

Trigger

```
CREATE TRIGGER score_log  
AFTER INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION  
log_score();
```

- This tells the database to

- Watch out for insertion on Score
- ~~Check for certain condition~~ *(we'll come back to this later)*
- Call the function log_score() after each insertion of a tuple
- There are other options, but that will be discussed later on

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

Triggers

- The Trigger Function

- RETURNS **TRIGGER** indicates that this is a trigger function
 - Can only RETURNS **TRIGGER**
- **NEW** refers to the new row inserted into Scores
 - Only accessible by trigger functions
- Other accessible data?
 - **TG_OP**
 - **TG_TABLE_NAME**
 - **OLD**
 - ...

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47



Table "Scores_Log"

<u>Name</u>	Date
Alice	2021-10-01
Bob	2021-10-09
Cathy	2021-10-12
David	2021-10-15

```
CREATE OR REPLACE FUNCTION log_score()  
RETURNS TRIGGER AS $$  
BEGIN  
  
    INSERT INTO Scores_Log  
        VALUES (NEW.Name, CURRENT_DATE);  
  
    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

Triggers

- Syntax

Trigger

```
CREATE TRIGGER <trigger_name>  
<timing> <event> ON <table>  
FOR EACH <granularity>  
EXECUTE FUNCTION <function_name>();
```

only applicable to this table

when triggered, this
trigger calls this function

NOTE

<timing>, <event> and
<granularity> are called
trigger options.

This is what we will discuss
next.

```
CREATE OR REPLACE FUNCTION <function_name>()  
RETURNS TRIGGER AS $$  
BEGIN
```

can only return TRIGGER

<code>

this code is then executed

```
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

Triggers Options

- Trigger Options

- Events:

- **INSERT ON** *table*
 - **DELETE ON** *table*
 - **UPDATE [OF column] ON** *table*

TG_OP {
'INSERT'
'DELETE'
'UPDATE'

- Timing:

- **AFTER** or **BEFORE** *(the triggering event)*
 - **INSTEAD OF** *(the triggering event on views)*

- Granularity:

- **FOR EACH ROW** *(modified)*
 - **FOR EACH STATEMENT** *(that performs the modification)*

Triggers Events

INSERT

DELETE

UPDATE

Triggers Event

- **Another Motivating Example**

Table "Scores"

<u>Name</u>	Mark
Alice	92
Bob	63
Cathy	58
David	47




Table "Scores_Log2"

<u>Name</u>	Op	Date
Alice	Insert	2021-10-01
Bob	Delete	2021-10-09
Cathy	Update	2021-10-12
David	Insert	2021-10-15

- Suppose we want to log when the marks are entered
 - This should be done automatically
 - The user of the database should not be bothered to write SQL statement to insert
 - This should be done *each time* an insertion occurs regardless of how it is done
 - We want to record the following data
 - The name of the student
 - The date of entry
 - The operation performed

Triggers Event

- The Trigger Function

```
CREATE OR REPLACE FUNCTION log_score2() RETURNS TRIGGER AS $$
BEGIN
    IF (TG_OP = 'INSERT') THEN
        INSERT INTO Scores_Log2 VALUES (NEW.Name, 'Insert', CURRENT_DATE);
        RETURN NEW;
    ELSIF (TG_OP = 'DELETE') THEN
        INSERT INTO Scores_Log2 VALUES (OLD.Name, 'Delete', CURRENT_DATE);
        RETURN OLD;
    ELSIF (TG_OP = 'UPDATE') THEN
        INSERT INTO Scores_Log2 VALUES (NEW.Name, 'Update', CURRENT_DATE);
        RETURN NEW;
    END IF;
END;
$$ LANGUAGE plpgsql;
```

Trigger Function

Question

Can we RETURN NULL like before?

Return value has effect, we will discuss later.

Triggers Event

- The Trigger Function

```
CREATE OR REPLACE FUNCTION log_score2() RETURNS TRIGGER AS $$
BEGIN
    IF (TG_OP = 'INSERT') THEN ...
    ELIF (TG_OP = 'DELETE') THEN ...
    ELIF (TG_OP = 'UPDATE') THEN ...
    END IF;
END;
$$ LANGUAGE plpgsql;
```

Trigger Function

- The Trigger

```
CREATE TRIGGER score_log2
AFTER INSERT OR DELETE OR UPDATE ON Scores
FOR EACH ROW EXECUTE FUNCTION log_score2();
```

Trigger

Triggers Timing

AFTER
BEFORE
INSTEAD OF


Triggers Timing

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$  
BEGIN  
    IF (NEW.Name = 'Adi') THEN  
        NEW.Mark := 100;  
    END IF;    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```



Trigger Function

```
CREATE TRIGGER adi_should_get_full_mark  
BEFORE INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

NOTE
NEW is a tuple (Adi, 100).
Hence, the result.

Trigger

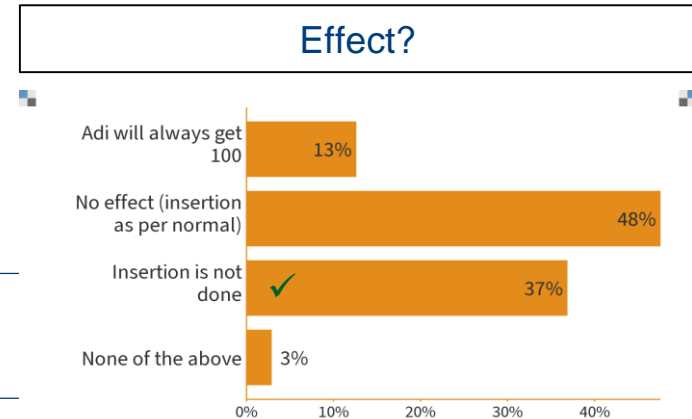
Triggers Return Value

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$  
BEGIN  
    IF (NEW.Name = 'Adi') THEN  
        NEW.Mark := 100;  
    END IF;    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```



Trigger Function

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark  
BEFORE INSERT ON Scores  
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

NOTE
NULL stops BEFORE trigger.
Hence, the result.

Trigger

Triggers Return Value

ERRATA

Here, OLD is initially NULL.
So it is the same as return NULL.

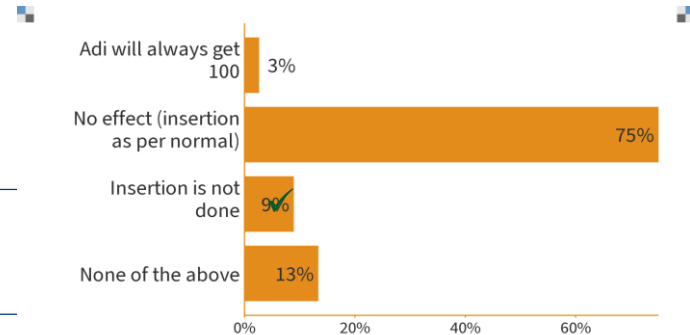
Table "Scores"

Name	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$
BEGIN
  IF (NEW.Name = 'Adi') THEN
    NEW.Mark := 100;
  END IF;  RETURN OLD;
END;
$$ LANGUAGE plpgsql;
```

Effect?



Trigger Function

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark
BEFORE INSERT ON Scores
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

NOTE

OLD is NULL for BEFORE trigger.
Hence, the result.

Trigger


Triggers Return Value

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$
BEGIN
    IF (NEW.Name = 'Adi') THEN
        OLD.Name := 'Adi';    OLD.Mark := 100;
    END IF;    RETURN OLD;
END;
$$ LANGUAGE plpgsql;
```



Trigger Function

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark
BEFORE INSERT ON Scores
FOR EACH ROW EXECUTE FUNCTION give_adi_full_mark();
```

NOTE
OLD is NULL for non-Adi.
Old is (Adi, 100) for Adi.
Hence, the result.

Trigger

Triggers Return Value

- **Transition Variables**

- **NEW** The modified row *after* the triggering event
- **OLD** The modified row *before* the triggering event
- ❖ Not all make sense all the time

- **Effect of Return Value**

NOTE
✗ is NULL

	NEW	OLD
INSERT	✓	✗
UPDATE	✓	✓
DELETE	✗	✓

	NULL tuple	non-NULL tuple t
BEFORE INSERT	No tuple inserted	Tuple t will be inserted
BEFORE UPDATE	No tuple updated	Tuple t will be the updated tuple
BEFORE DELETE	No deletion performed	Deletion proceeds as normal
AFTER INSERT	Does not matter <i>It's done already, cannot be changed now</i>	
AFTER UPDATE		
AFTER DELETE		

Triggers Timing

- INSTEAD OF Trigger
 - This kind of trigger can only be defined on **VIEWS**
- First Off...
 - What is a **VIEW**?
 - A “*virtual*” table defined by a query to compute the view
 - Can be used in queries just like a regular table
 - Why use a **VIEW**?
 - Hide data and/or complexity from users
 - Logical data independence
 - Captures common query to be reused
 - ❖ Real database applications use **tons** of VIEW

Triggers Timing

- **INSTEAD OF Trigger**
 - This kind of trigger can only be defined on **VIEWS**
- **Secondly ...**
 - Why do we want to modify a **VIEW**?
 - User sees it as a table and they can modify table
 - Does it even make sense?
 - In most cases, we cannot modify **VIEW** directly
 - But we can modify the underlying table

Triggers Timing

- **INSTEAD OF Trigger**

- This kind of trigger can only be defined on **VIEWS**

- **Modifying a VIEW**

- Simple case

- `CREATE VIEW Students AS
SELECT Name FROM Scores;`

- `DELETE FROM Students WHERE Name = 'Alice';`
translates to
`DELETE FROM Scores WHERE Name = 'Alice';`

Triggers Timing

- INSTEAD OF Trigger

- This kind of trigger can only be defined on **VIEWS**

- Modifying a VIEW

- Impossible case

- ```
CREATE VIEW A_Students AS
SELECT Name, Mark FROM Scores WHERE Mark >= 70;
```
- ```
INSERT INTO A_Students VALUES ('Bryan', 69);
```

insertion does not affect the VIEW, only the underlying table

Triggers Timing

- INSTEAD OF Trigger

- This kind of trigger can only be defined on **VIEWS**

- Modifying a VIEW

- Too many possible case
 - `CREATE VIEW Top_Marks(Mark) AS
SELECT MAX(Mark) FROM Scores;`
 - `UPDATE Top_Marks SET Mark = 80;`
what if more than one people have top marks?

Triggers Timing

- **INSTEAD OF Trigger**
 - This kind of trigger can only be defined on **VIEWS**
- **OK, so what to do here?**
 - Let's say we want to modify all people with the maximum mark
 - We modify the underlying table instead
 - Use INSTEAD OF trigger

```
CREATE TRIGGER modify_top_mark  
INSTEAD OF UPDATE ON Top_Marks  
FOR EACH ROW EXECUTE FUNCTION update_top_mark();
```

Trigger

Triggers Timing

Table "Scores"

<u>Name</u>	Mark
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION update_top_mark() RETURNS TRIGGER AS $$  
BEGIN  
    UPDATE Scores SET Mark = NEW.Mark WHERE Mark = OLD.Mark;  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

Return Value

- **NULL**
 - Ignore all operations on current row
- **Non-NULL**
 - Signals the database to proceed as normal

- The Trigger

```
CREATE TRIGGER modify_top_mark  
INSTEAD OF UPDATE ON Top_Marks  
FOR EACH ROW EXECUTE FUNCTION update_top_mark();
```

Trigger

Triggers Granularity

FOR EACH ROW

FOR EACH STATEMENT

Triggers Granularity

- **Row-Level Trigger** *(FOR EACH ROW)*
 - Executes the trigger function for every tuple encountered
- **Statement-Level Trigger** *(FOR EACH STATEMENT)*
 - Executes the trigger function only once
 - Why would we want to do this?
 - If we already prevent an operation, then just once is enough
 - There are other contextual data available, but we will not discuss this

```
CREATE TRIGGER modify_top_mark  
INSTEAD OF UPDATE ON Top_Marks  
FOR EACH ROW EXECUTE FUNCTION update_top_mark();
```

Triggers Granularity

Table "Scores_Log"

<u>Name</u>	Date
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION show_warning() RETURNS TRIGGER AS $$  
BEGIN  
    RAISE NOTICE 'You are not supposed to delete from log...';  
    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

- The Trigger

```
CREATE TRIGGER warn_delete  
BEFORE DELETE ON Scores_Log  
FOR EACH STATEMENT EXECUTE FUNCTION show_warning();
```

Trigger

Effect?

- **RAISE NOTICE**
 - Database will give you prompt whenever a deletion is attempted
- **RETURN NULL**
 - Unfortunately, still perform deletion

Triggers Granularity

Table "Scores_Log"

<u>Name</u>	Date
...	...

- The Trigger Function

```
CREATE OR REPLACE FUNCTION show_warning() RETURNS TRIGGER AS $$  
BEGIN  
    RAISE EXCEPTION 'You are not supposed to delete from log...';  
    RETURN NULL;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Function

- Statement-Level Trigger *(FOR EACH STATEMENT)*

- Ignores the values returned by the trigger functions
- RETURN NULL would not make the database omit the subsequent operation
- What to do?
 - RAISE EXCEPTION

Triggers Granularity

- **Granularity and Timing**

- **INSTEAD OF** is only allowed on row-level granularity
- **AFTER** or **BEFORE** are allowed on both row-level as well as statement-level granularity

Granularity

Timing	Row-Level	Statement-Level
AFTER	Tables	Tables and View
BEFORE	Tables	Tables and View
INSTEAD OF	Views	-

Triggers Condition

Event

Trigger

Condition

... we are here ...

Action

Trigger Function

Triggers Condition

Table "Scores"

<u>Name</u>	Mark
Adi	100

- The Trigger Function

```
CREATE OR REPLACE FUNCTION give_adi_full_mark() RETURNS TRIGGER AS $$
BEGIN
  IF (NEW.Name = 'Adi') THEN
    NEW.Mark := 100;
  END IF;
  RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

Trigger Function

Observation

Trigger function only cares about the case when name is *Adi*

We can move this condition to the trigger definition

- The Trigger

```
CREATE TRIGGER adi_should_get_full_mark
BEFORE INSERT ON Scores
FOR EACH ROW WHEN (NEW.Name = 'Adi')
EXECUTE FUNCTION give_adi_full_mark();
```

Trigger

Triggers Condition

- What Can We Use?


- In general, the condition in **WHEN()** could be more complicated
- Subject to the following requirements:
 - **NO SELECT** in **WHEN()**
 - **NO OLD** in **WHEN()** for **INSERT**
 - **NO NEW** in **WHEN()** for **DELETE**
 - **NO WHEN()** for **INSTEAD OF**

```
CREATE TRIGGER adi_should_get_full_mark
BEFORE INSERT ON Scores
FOR EACH ROW WHEN (NEW.Name = 'Adi')
EXECUTE FUNCTION give_adi_full_mark();
```


Triggers Condition

- **Deferred Trigger**

- Triggers happen at the end of either statement or at the end of transaction
- **Operation** consisting of multiple statements may leave the database in an intermediate *inconsistent* state
- In such cases, we want the trigger to check consistency constraint only at the end of a transaction
- ❖ This is called a *deferred* trigger



Triggers Condition


- **Deferred Trigger**

- Example:
 - Customer may have multiple account
 - Total balance for all accounts must be at least 150

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- Task: Transfer money from one account to another



Triggers Condition

- **Deferred Trigger**

- Naïve approach:

1. Deduct amount from account 1
2. Add amount to account 2


Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- Solution:

1. Put the two update into one transaction
2. Defer trigger to check only at the end of transaction

May violate constraint!



Triggers Condition

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- Deferred Trigger

```
CREATE CONSTRAINT TRIGGER balance_check  
AFTER INSERT OR UPDATE OR DELETE ON Account  
DEFERRABLE INITIALLY DEFERRED  
FOR EACH ROW EXECUTE FUNCTION check_balance();
```

- **CONSTRAINT** and **DEFERRABLE** *together* indicate that the trigger can be deferred
- **INITIALLY DEFERRED** indicates that by default, the trigger is deferred
 - In other words, only check at the end of transaction
 - Other option is **INITIALLY IMMEDIATE**
 - *i.e.*, the trigger is not deferred by default
- Deferred triggers only work with **AFTER** and **FOR EACH ROW**



Triggers Condition

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- **Deferred Trigger**

```
CREATE CONSTRAINT TRIGGER balance_check  
AFTER INSERT OR UPDATE OR DELETE ON Account  
DEFERRABLE INITIALLY DEFERRED  
FOR EACH ROW EXECUTE FUNCTION check_balance();
```

- Now we can do the following

```
BEGIN TRANSACTION;  
  
UPDATE Account SET Balance = Balance - Amount WHERE AID = Account1;  
UPDATE Account SET Balance = Balance + Amount WHERE AID = Account2;  
  
COMMIT;
```

Triggers Condition

Table "Account"

<u>AID</u>	Name	Balance
1	Alice	100
2	Alice	100

- **Deferred Trigger**

```
CREATE CONSTRAINT TRIGGER balance_check  
AFTER INSERT OR UPDATE OR DELETE ON Account  
DEFERRABLE INITIALLY IMMEDIATE  
FOR EACH ROW EXECUTE FUNCTION check_balance();
```

- What if the trigger is INITIALLY IMMEDIATE?

- Change it on the fly

```
BEGIN TRANSACTION;  
SET CONSTRAINTS balance_check DEFERRED;  
UPDATE Account SET Balance = Balance - Amount WHERE AID = Account1;  
UPDATE Account SET Balance = Balance + Amount WHERE AID = Account2;  
  
COMMIT;
```

Final Note on Triggers

- **Multiple Triggers**

- There can be multiple triggers defined for the same event on the same table
 - There need to be an *order of activation*
 - BEFORE statement-level triggers
 - BEFORE row-level triggers
 - AFTER row-level triggers
 - AFTER statement-level triggers
 - Within each category, triggers are activated in *alphabetic* order
 - If BEFORE row-level trigger returns NULL, then subsequent triggers on the same row are omitted

- **Universality of Triggers?**

- Our discussions are based on PostgreSQL syntax and implementation only

Questions?